




# Pandas



- Let's quickly review our ML Pathway...




# ML Pathway





# ML Pathway





# ML Pathway



**Real  
World**



**Collect &  
Store  
Data**



**Clean &  
Organize  
Data**



**Exploratory  
Data  
Analysis**



 **pandas**



# Pandas

- Pandas is a library for Data Analysis.
- Extremely powerful table (DataFrame) system built off of NumPy.
- Fantastic documentation:
  - <https://pandas.pydata.org/docs/>





# Pandas

- What can we do with Pandas?
  - Tools for reading and writing data between many formats.
  - Intelligently grab data based on indexing, logic, subsetting, and more.
  - Handle missing data.
  - Adjust and restructure data.



# Pandas - Section Overview

- Series and DataFrames
- Conditional Filtering and Useful Methods
- Missing Data
- Group By Operations
- Combining DataFrames
- Text Methods and Time Methods
- Inputs and Outputs





# Let's get started!



# Series

- A Series is a data structure in Pandas that holds an array of information along with a named index.
- The named index differentiates this from a simple NumPy array.
- **Formal Definition:** One-dimensional ndarray with axis labels

- NumPy array has numeric index

<b>0</b>	<b>1776</b>
<b>1</b>	<b>1867</b>
<b>2</b>	<b>1821</b>

- NumPy array has numeric index

Index	Data
0	1776
1	1867
2	1821

- Pandas Series adds on a labeled index

Labeled Index	Data
USA	1776
CANADA	1867
MEXICO	1821

- Data is still numerically organized

<b>Numeric Index</b>	<b>Labeled Index</b>	<b>Data</b>
<b>0</b>	<b>USA</b>	<b>1776</b>
<b>1</b>	<b>CANADA</b>	<b>1867</b>
<b>2</b>	<b>MEXICO</b>	<b>1821</b>



- Let's explore the various ways to create a Pandas Series object.
- We'll also learn about some key properties and operations.
- Later on we will learn how to combine Series with a shared index to create a tabular data structure called a DataFrame.





# Series

PART TWO



# DataFrames

Part One



- A DataFrame is a table of columns and rows in pandas that we can easily restructure and filter.
- **Formal Definition:** A group of Pandas Series objects that *share* the same index.

- Example of a Series

Index	Year
USA	1776
CANADA	1867
MEXICO	1821

- Example of Series with Same Index

Index	Year
USA	1776
CANADA	1867
MEXICO	1821

Index	Pop
USA	328
CANADA	38
MEXICO	126

Index	GDP
USA	20.5
CANADA	1.7
MEXICO	1.22

- Example of Series with Same Index

Index	Year
USA	1776
CANADA	1867
MEXICO	1821

Index	Pop
USA	328
CANADA	38
MEXICO	126

Index	GDP
USA	20.5
CANADA	1.7
MEXICO	1.22

- DataFrame

Index	Year	Pop	GDP
USA	1776	328	20.5
CANADA	1867	38	1.7
MEXICO	1821	126	1.22



- DataFrame is the main Pandas object we will work with and it is **extremely** useful!
- This series covers first the “basics”
  - Create a DataFrame
  - Grab a column or multiple columns
  - Grab a row or multiple rows
  - Insert a new column or new row



- ***Quick Note: Each video lecture in this DataFrames series refers to the same 01-DataFrames.ipynb notebook!***



# DataFrames

Part Two



# DataFrames

Part Three



# DataFrames

Part Four



# Conditional Filtering



- Typically in data analysis our datasets are large enough that we don't filter based on position, but instead based on a **condition**.
- Conditional Filtering allows us to select **rows** based a condition on a column.
- This leads to a discussion on organizing our data...

- Organizing Data

Index	Year	Pop	GDP
USA	1776	328	20.5
CANADA	1867	38	1.7
MEXICO	1821	126	1.22

- Columns are Features

Index	Year	Pop	GDP
USA	1776	328	20.5
CANADA	1867	38	1.7
MEXICO	1821	126	1.22



- Rows are instances of data

Index	Year	Pop	GDP
USA	1776	328	20.5
CANADA	1867	38	1.7
MEXICO	1821	126	1.22

- This format is required for ML later on!

Index	Year	Pop	GDP
USA	1776	328	20.5
CANADA	1867	38	1.7
MEXICO	1821	126	1.22

- This allows to directly answer questions

Index	Year	Pop	GDP
USA	1776	328	20.5
CANADA	1867	38	1.7
MEXICO	1821	126	1.22

- What countries have Pop greater than X?

Index	Year	Pop	GDP
USA	1776	328	20.5
CANADA	1867	38	1.7
MEXICO	1821	126	1.22

- What countries have Pop greater than 50?

Index	Year	Pop	GDP
USA	1776	328	20.5
CANADA	1867	38	1.7
MEXICO	1821	126	1.22



- `df["Pop"]`

Index	Year	Pop	GDP
USA	1776	328	20.5
CANADA	1867	38	1.7
MEXICO	1821	126	1.22

- `df["Pop"] > 50`

Index	Year	Pop	GDP
USA	1776	328	20.5
CANADA	1867	38	1.7
MEXICO	1821	126	1.22

- `df["Pop"] > 50`

Index	Year	Pop	GDP
USA	1776	328	20.5
CANADA	1867	38	1.7
MEXICO	1821	126	1.22



- `df["Pop"] > 50`

Index	Year	Pop	GDP
USA	1776	True	20.5
CANADA	1867	False	1.7
MEXICO	1821	True	1.22

- `df[df["Pop"] > 50]`

Index	Year	Pop	GDP
USA	1776	True	20.5
CANADA	1867	False	1.7
MEXICO	1821	True	1.22

- `df[df["Pop"] > 50]`

Index	Year	Pop	GDP
USA	1776	True	20.5
MEXICO	1821	True	1.22



# Pandas

- Conditional Filtering:
  - Filter by single condition
  - Filter by multiple conditions
  - Check against multiple possible values



# Useful Methods

PART ONE - APPLY METHODS



- We now understand the basics of how to grab and filter data from a Series or DataFrame in pandas.
- We are now going to cover a wide variety of method calls available in Pandas.
- This will be part of a series of lectures since there are quite a few methods to cover.



# Pandas

- For your convenience, the lecture notebook for this series has a list at the top with links that take you directly to the relevant section of the notebook for a topic.



- While pandas has many built in methods, we can use the `.apply()` method call to apply any custom python function of our own to every row in a Series.
- We can use either one or multiple columns as input, let's explore this in the notebook!





# Useful Methods

PART TWO - APPLY WITH MULTIPLE COLUMNS



# Useful Methods

PART THREE - DESCRIBING AND SORTING



# Useful Methods

PART THREE - METHOD CALLS



# Missing Data

PART ONE - OVERVIEW

- Real world data will often be missing data for a variety of reasons.
- Many machine learning models and statistical methods can not work with missing data points, in which case we need to decide what to do with the missing data.



- When reading in missing values, pandas will display them as **NaN** values.
- There are also newer specialized null pandas values such as **pd.NaT** to imply the value missing should be a timestamp.



# Pandas

- Options for Missing Data
  - Keep it
  - Remove it
  - Replace it
- *Note, there is never 100% correct approach that applies to all circumstances, it all depends on the exact situation you encounter!*



# Pandas

- Keeping the missing data
  - PROS:
    - Easiest to do
    - Does not manipulate or change the true data
  - CONS:
    - Many methods do not support NaN
    - Often there are reasonable guesses





# Pandas

- Dropping or Removing the missing data
  - PROS:
    - Easy to do.
    - Can be based on rules.
  - CONS:
    - Potential to lose a lot of data or useful information.
    - Limits trained models for future data.



# Pandas

- Removing or Dropping missing data
  - Dropping a Row
    - Makes sense when a lot of info is missing

	Year	Pop	GDP	Area
USA	1776	NAN	NAN	NAN
CANADA	1867	38	1.7	3.86
MEXICO	1821	126	1.22	0.76



# Pandas

- Removing or Dropping missing data
  - Dropping a Row
    - Clearly this data point as a row should probably be dropped

	Year	Pop	GDP	Area
USA	1776	NAN	NAN	NAN
CANADA	1867	38	1.7	3.86
MEXICO	1821	126	1.22	0.76



# Pandas

- Removing or Dropping missing data
  - Dropping a Row
    - Often a good idea to calculate a percentage of what data is dropped

	Year	Pop	GDP	Area
USA	1776	NAN	NAN	NAN
CANADA	1867	38	1.7	3.86
MEXICO	1821	126	1.22	0.76



# Pandas

- Removing or Dropping missing data
  - Dropping a Column
    - Good choice if every row is missing that particular feature

	Year	Pop	GDP	Area
USA	1776	328	20.5	NAN
CANADA	1867	38	1.7	NAN
MEXICO	1821	126	1.22	0.76



# Pandas

- Filling in the missing data
  - PROS:
    - Potential to save a lot of data for use in training a model
  - CONS:
    - Hardest to do and somewhat arbitrary
    - Potential to lead to false conclusions



# Pandas

- Filling in missing data
  - Fill with same value
    - Good choice if NaN was a placeholder

	Year	Pop	GDP	Carriers
USA	1776	328	20.5	11
CANADA	1867	38	1.7	NAN
MEXICO	1821	126	1.22	NAN



# Pandas

- Filling in missing data
  - Fill with same value
    - Good choice if NaN was a placeholder

	Year	Pop	GDP	Carriers
USA	1776	328	20.5	11
CANADA	1867	38	1.7	NAN
MEXICO	1821	126	1.22	NAN





# Pandas

- Filling in missing data
  - Fill with same value
    - Here NAN can be filled in with zero

	Year	Pop	GDP	Carriers
USA	1776	328	20.5	11
CANADA	1867	38	1.7	0
MEXICO	1821	126	1.22	0

- Filling in missing data
  - Fill with interpolated or estimated value
    - Much harder and requires reasonable assumptions

	Year	Pop	GDP	Perct
USA	1776	328	20.5	75%
CANADA	1867	38	1.7	NAN
MEXICO	1821	126	1.22	25%



# Pandas

- Filling in missing data
  - Fill with interpolated or estimated value
    - Much harder and requires reasonable assumptions

	Year	Pop	GDP	Perct
USA	1776	328	20.5	75%
CANADA	1867	38	1.7	50%
MEXICO	1821	126	1.22	25%





# Pandas

- Let's explore the code syntax in pandas for dealing with missing values.
- Later on in the course we will have a deeper discussion on trying to decide between keep, remove, and replace options.



# Missing Data

PART TWO - PANDAS



# Groupby Operations



# Pandas

- A `groupby()` operation allows us to examine data on a **per category** basis.
- Let's explore what this looks like in pandas...



Category	Data Value
A	10
A	5
B	2
B	4
C	12
C	6



Category	Data Value
A	10
A	5
B	2
B	4
C	12
C	6

We need to choose a **categorical** column to call with **groupby()**.

Categorical columns are non-continuous.

Keep in mind, they can still be numerical, such as cabin class categories on a ship (e.g. Class 1, Class 2, Class 3)

Let's now see what happens with a `.groupby()` call combined with an aggregate function call.

Category	Data Value
A	10
A	5
B	2
B	4
C	12
C	6



# Pandas


Category	Data Value
A	10
A	5
B	2
B	4
C	12
C	6



A	10
A	5



B	2
B	4



C	12
C	6



# Pandas

Category	Data Value
A	10
A	5
B	2
B	4
C	12
C	6

A	10
A	5

B	2
B	4

C	12
C	6

Aggregate Function  
**.sum()**

Category	Result
A	15
B	6
C	18



# Pandas

Category	Data Value
A	10
A	5
B	2
B	4
C	12
C	6

A	10
A	5

B	2
B	4

C	12
C	6

Aggregate Function  
**.mean()**

Category	Result
A	7.5
B	3
C	9



# Pandas

Category	Data Value
A	10
A	5
B	2
B	4
C	12
C	6

A	10
A	5

B	2
B	4

C	12
C	6

Aggregate Function  
**.count()**

Category	Result
A	2
B	2
C	2



- Note that in pandas calling **groupby()** by itself creates a “*lazy*” groupby object waiting to be evaluated by an aggregate method call.
- Let's explore this further in pandas!



# Groupby Operations

MULTI-LEVEL INDEX CONTINUED...





# Combining DataFrames

Concatenation

- Often the data you need exists in two separate sources, fortunately, Pandas makes it easy to combine these together.
- The simplest combination is if both sources are already in the same format, then a **concatenation** through the **pd.concat()** call is all that is needed.

- Concatenation is simply “pasting” the two DataFrames together, by columns:

	Year	Pop
USA	1776	328
CANADA	1867	38
MEXICO	1821	126



	GDP	Perct
USA	20.5	75%
CANADA	1.7	NAN
MEXICO	1.22	25%

- Concatenation is simply “pasting” the two DataFrames together, by columns:

	Year	Pop	GDP	Perct
USA	1776	328	20.5	75%
CANADA	1867	38	1.7	NAN
MEXICO	1821	126	1.22	25%

- Concatenation is simply “pasting” the two DataFrames together, by rows:

	Year	Pop	GDP
USA	1776	328	20.5
CANADA	1867	38	1.7



	Year	Pop	GDP
MEXICO	1821	126	1.22
BRAZIL	1822	209	1.86

- Concatenation is simply “pasting” the two DataFrames together, by rows:

	Year	Pop	GDP
USA	1776	328	20.5
CANADA	1867	38	1.7
BRAZIL	1822	209	1.86
MEXICO	1821	126	1.22



- Pandas will also automatically fill NaN where necessary.
- Let's explore some examples in Pandas!



# Combining DataFrames

“Inner” Merge





# Pandas

- Often DataFrames are not in the exact same order or format, meaning we can not simply concatenate them together.
- In this case, we need to **merge** the DataFrames.
- This is analogous to a JOIN command in SQL.



# Pandas

- The `.merge()` method takes in a key argument labeled **how**
- There are 3 main ways of merging tables together using the **how** parameter:
  - Inner
  - Outer
  - Left or Right



- The main idea behind the argument is to decide **how** to deal with information only present in one of the joined tables.



# Pandas

- Let's imagine a simple example.
- Our company is holding a conference for people in the movie rental industry.
- We'll have people register online beforehand and then login the day of the conference.

- After the conference we have these tables

REGISTRATIONS	
reg_id	name
1	Andrew
2	Bob
3	Charlie
4	David

LOGINS	
log_id	name
1	Xavier
2	Andrew
3	Yolanda
4	Bob

- The respective id columns indicate what order they registered or logged in on site.

REGISTRATIONS	
reg_id	name
1	Andrew
2	Bob
3	Charlie
4	David

LOGINS	
log_id	name
1	Xavier
2	Andrew
3	Yolanda
4	Bob

- For the sake of simplicity, we will assume the names are unique.

REGISTRATIONS	
reg_id	name
1	Andrew
2	Bob
3	Charlie
4	David

LOGINS	
log_id	name
1	Xavier
2	Andrew
3	Yolanda
4	Bob

- (e.g. There is only one person in the company named “Andrew”)

REGISTRATIONS	
reg_id	name
1	Andrew
2	Bob
3	Charlie
4	David

LOGINS	
log_id	name
1	Xavier
2	Andrew
3	Yolanda
4	Bob



- To help you keep track, Registrations names' first letters go A,B,C,D

REGISTRATIONS	
reg_id	name
1	Andrew
2	Bob
3	Charlie
4	David

LOGINS	
log_id	name
1	Xavier
2	Andrew
3	Yolanda
4	Bob

- First we need to decide **on** what column to merge together.

REGISTRATIONS	
reg_id	name
1	Andrew
2	Bob
3	Charlie
4	David

LOGINS	
log_id	name
1	Xavier
2	Andrew
3	Yolanda
4	Bob

- The **on** column should be a *primary* identifier, meaning unique per row.

REGISTRATIONS	
reg_id	name
1	Andrew
2	Bob
3	Charlie
4	David

LOGINS	
log_id	name
1	Xavier
2	Andrew
3	Yolanda
4	Bob

- The **on** column should also be present in both tables being merged.

REGISTRATIONS	
reg_id	name
1	Andrew
2	Bob
3	Charlie
4	David

LOGINS	
log_id	name
1	Xavier
2	Andrew
3	Yolanda
4	Bob

- Since we assume names are unique here, will we merge **on= "name"**.

REGISTRATIONS	
reg_id	name
1	Andrew
2	Bob
3	Charlie
4	David

LOGINS	
log_id	name
1	Xavier
2	Andrew
3	Yolanda
4	Bob

- Next we need to decide **how** to merge the tables **on** the **name** column.

REGISTRATIONS	
reg_id	name
1	Andrew
2	Bob
3	Charlie
4	David

LOGINS	
log_id	name
1	Xavier
2	Andrew
3	Yolanda
4	Bob

- With **how="inner"** the result will be the set of records that match in both tables.

REGISTRATIONS	
reg_id	name
1	Andrew
2	Bob
3	Charlie
4	David

LOGINS	
log_id	name
1	Xavier
2	Andrew
3	Yolanda
4	Bob

- With **how= “inner”** the result will be the set of records that **match in both** tables.

REGISTRATIONS	
reg_id	name
1	Andrew
2	Bob
3	Charlie
4	David


LOGINS	
log_id	name
1	Xavier
2	Andrew
3	Yolanda
4	Bob



## Merges are often shown as a Venn diagram

```
pd.merge(registrations,logins,how='inner',on='name')
```

REGISTRATIONS	
reg_id	name
1	Andrew
2	Bob
3	Charlie
4	David




LOGINS	
log_id	name
1	Xavier
2	Andrew
3	Yolanda
4	Bob



```
pd.merge(registrations,logins,how='inner',on='name')
```

REGISTRATIONS	
reg_id	name
1	Andrew
2	Bob
3	Charlie
4	David




LOGINS	
log_id	name
1	Xavier
2	Andrew
3	Yolanda
4	Bob



```
pd.merge(registrations,logins,how='inner',on='name')
```

REGISTRATIONS	
reg_id	name
1	Andrew
2	Bob
3	Charlie
4	David



LOGINS	
log_id	name
1	Xavier
2	Andrew
3	Yolanda
4	Bob



```
pd.merge(registrations,logins,how='inner',on='name')
```

REGISTRATIONS	
reg_id	name
1	Andrew
2	Bob
3	Charlie
4	David

RESULTS		
reg_id	name	log_id
1	Andrew	2
2	Bob	4

LOGINS	
log_id	name
1	Xavier
2	Andrew
3	Yolanda
4	Bob



- Let's quickly explore this in pandas!



# Combining DataFrames

“left” and “right” merge

- Now that we understand an “inner” merge, let’s explore “left” versus “right” merge conditions.
- Note! Order of the tables passed in as arguments does matter here!

Let's explore an **how= "left"** condition with our two example tables.

REGISTRATIONS	
reg_id	name
1	Andrew
2	Bob
3	Charlie
4	David

LOGINS	
log_id	name
1	Xavier
2	Andrew
3	Yolanda
4	Bob



Note: Registrations is the left table, logins will be the right table


REGISTRATIONS	
reg_id	name
1	Andrew
2	Bob
3	Charlie
4	David

LOGINS	
log_id	name
1	Xavier
2	Andrew
3	Yolanda
4	Bob



```
pd.merge(registrations,logins,how='left',on='name')
```

REGISTRATIONS	
reg_id	name
1	Andrew
2	Bob
3	Charlie
4	David



LOGINS	
log_id	name
1	Xavier
2	Andrew
3	Yolanda
4	Bob



```
pd.merge(registrations,logins,how='left',on='name')
```

REGISTRATIONS	
reg_id	name
1	Andrew
2	Bob
3	Charlie
4	David

RESULTS		
reg_id	name	log_id
1	Andrew	2
2	Bob	4
3	Charlie	NaN
4	David	NaN

LOGINS	
log_id	name
1	Xavier
2	Andrew
3	Yolanda
4	Bob



```
pd.merge(registrations,logins,how='left',on='name')
```

REGISTRATIONS	
reg_id	name
1	Andrew
2	Bob
3	Charlie
4	David

RESULTS		
reg_id	name	log_id
1	Andrew	2
2	Bob	4
3	Charlie	NaN
4	David	NaN

LOGINS	
log_id	name
1	Xavier
2	Andrew
3	Yolanda
4	Bob



- Now let's see what happens in a how="right" situation.



```
pd.merge(registrations,logins,how='right',on='name')
```

REGISTRATIONS	
reg_id	name
1	Andrew
2	Bob
3	Charlie
4	David

RESULTS		
reg_id	name	log_id
1	Andrew	2
2	Bob	4
NaN	Xavier	1
NaN	Yolanda	3

LOGINS	
log_id	name
1	Xavier
2	Andrew
3	Yolanda
4	Bob



- Let's explore this further in pandas!



# Combining DataFrames

“outer” merge



- Setting **how= “outer”** allows us to include everything present in both tables.

- Recall we match Andrew and Bob in both tables

REGISTRATIONS	
reg_id	name
1	Andrew
2	Bob
3	Charlie
4	David

LOGINS	
log_id	name
1	Xavier
2	Andrew
3	Yolanda
4	Bob

- But we have names that only appear in one table!

REGISTRATIONS	
reg_id	name
1	Andrew
2	Bob
3	Charlie
4	David

LOGINS	
log_id	name
1	Xavier
2	Andrew
3	Yolanda
4	Bob

- We can use **how= “outer”** to make sure we grab all names from both tables.


REGISTRATIONS	
reg_id	name
1	Andrew
2	Bob
3	Charlie
4	David

LOGINS	
log_id	name
1	Xavier
2	Andrew
3	Yolanda
4	Bob



```
pd.merge(registrations,logins,how='outer',on='name')
```

REGISTRATIONS	
reg_id	name
1	Andrew
2	Bob
3	Charlie
4	David



LOGINS	
log_id	name
1	Xavier
2	Andrew
3	Yolanda
4	Bob



```
pd.merge(registrations,logins,how='outer',on='name')
```

REGISTRATIONS	
reg_id	name
1	Andrew
2	Bob
3	Charlie
4	David

RESULTS		
reg_id	name	log_id
1	Andrew	2
2	Bob	4
3	Charlie	NaN
4	David	NaN
NaN	Xavier	1
NaN	Yolanda	3

LOGINS	
log_id	name
1	Xavier
2	Andrew
3	Yolanda
4	Bob



- Let's quickly explore this result in pandas!



# Combining DataFrames

Joining on Index and Different Key Names





# Text Methods



# Pandas

- Often text data needs to be cleaned or manipulated for processing.
- While we can always use a custom `apply()` function for these tasks, pandas comes with many built-in string method calls.
- Let's learn how to use them!



# Time Methods



# Pandas

- Basic Python has a **datetime** object containing date and time information.
- Pandas allows us to easily extract information from a datetime object to use feature engineering.



# Pandas

- For example, we may have recent timestamped sales data.
- Pandas will allow us to extract information from the timestamp, such as:
  - Day of the Week
  - Weekend vs Weekday
  - AM vs PM



# Data Input and Output

CSV Files



# Pandas

- Pandas can read in data from a wide variety of sources and has excellent online documentation!
- In this series of lectures we will cover some of the most popular ways to read in datasets.



# Pandas

- Note!
  - You need to know the **exact** directory location and correct file name.
  - You may need passwords or permissions for certain data inputs (e.g. a SQL database password).





# Pandas

- Final Note:
  - It's almost impossible for us to help with datasets outside the course, since they could be incorrectly formatted, in the wrong location, or have a different name.



# Pandas

- Video Lectures:
  - CSV Files
  - HTML Tables
  - Excel Files
  - SQL Databases



# Data Input and Output

HTML Tables



# Pandas

- Websites display tabular information through the use of HTML tables tags:
  - **<table>**
- Pandas has the ability to automatically convert these HTML tables into a DataFrame.

- *Important Notes!*
  - Not every table in a website is available through HTML tables.
  - Some websites may block your computer from scraping the HTML of the site through pandas.
  - It may be more efficient to use an API.



- Let's work through an example of grabbing all the tables from a Wikipedia Article and then cleaning and organizing the information to get a DataFrame.
- Output to an HTML table is also very useful to display tables on a website!



# Data Input and Output

Excel Files



# Pandas

- Pandas can read and write to Excel files.
- *Important Note!*
  - Pandas can only read and write in raw data, it is not able to read in macros, visualizations, or formulas created inside of spreadsheets.





- Pandas treats an Excel Workbook as a dictionary, with the key being the sheet name and the value being the DataFrame representing the sheet itself.
- *Note! Using pandas with Excel requires additional libraries!*
- Let's explore how this works!



# Data Input and Output

SQL



# Pandas

- Pandas can read and write to various SQL engines through the use of a driver and the **sqlalchemy** python library.
- So how does this work?



# Pandas

- Step 1:
  - Figure out what SQL Engine you are connecting to, for just a few examples:
    - PostgreSQL
    - MySQL
    - MS SQL Server



# Pandas

- Step 2:
  - Install the appropriate Python driver library (*Most likely requires a Google Search*):
    - PostgreSQL - *psycopg2*
    - MySQL - *pymysql*
    - MS SQL Server - *pyodbc*



# Pandas

- Step 3:
  - Use the sqlalchemy library to connect to your SQL database with the driver:
    - [docs.sqlalchemy.org/en/13/dialects/index.html](https://docs.sqlalchemy.org/en/13/dialects/index.html)



# Pandas

- Step 4:
  - Use the sqlalchemy driver connection with pandas read\_sql method
  - Pandas can read in entire tables as a DataFrame or actual parse a SQL query through the connection:
    - `SELECT * FROM table;`



# Pandas

- *Important Note!*
  - It's almost impossible for us to help with your specific work databases outside of the course material, since it requires knowledge of your permissions, database names and locations, and password information!





# Pandas

- *Important Note!*
  - Use your skills in information lookup to easily find many online resources regarding examples for all of the major SQL engines, for example:
    - Google Search: *Oracle SQL + pandas*



# Pandas

- For our example, we'll use SQLite since it comes with Python and we can easily create a temporary database inside of your RAM.



# Pivot Tables



# Pandas

- Pivot tables allow you to reorganize data, refactoring cells based on columns and a new index.
- This is best shown visually...



# Pandas

- A DataFrame with repeated values can be pivoted for a reorganization and clarity

df

	foo	bar	baz	zoo
0	one	A	1	x
1	one	B	2	y
2	one	C	3	z
3	two	A	4	q
4	two	B	5	w
5	two	C	6	t



```
df.pivot(index='foo',  
          columns='bar',  
          values='baz')
```

bar	A	B	C
foo			
one	1	2	3
two	4	5	6



# Pandas

- We choose columns to define the new index, columns, and values.

df

	foo	bar	baz	zoo
0	one	A	1	x
1	one	B	2	y
2	one	C	3	z
3	two	A	4	q
4	two	B	5	w
5	two	C	6	t



```
df.pivot(index='foo',  
          columns='bar',  
          values='baz')
```


bar	A	B	C
foo			
one	1	2	3
two	4	5	6

- Notice how the choices for index and column should have repeated values.

df

```
df.pivot(index='foo',  
         columns='bar',  
         values='baz')
```

	foo	bar	baz	zoo
0	one	A	1	x
1	one	B	2	y
2	one	C	3	z
3	two	A	4	q
4	two	B	5	w
5	two	C	6	t



bar	A	B	C
foo			
one	1	2	3
two	4	5	6

- Also notice how all the information from the **zoo** column is now discarded.

df

```
df.pivot(index='foo',  
          columns='bar',  
          values='baz')
```

	foo	bar	baz	zoo
0	one	A	1	x
1	one	B	2	y
2	one	C	3	z
3	two	A	4	q
4	two	B	5	w
5	two	C	6	t

➔

bar	A	B	C
foo			
one	1	2	3
two	4	5	6





# Pandas

- No new information is shown, it is merely reorganized.

df

	foo	bar	baz	zoo
0	one	A	1	x
1	one	B	2	y
2	one	C	3	z
3	two	A	4	q
4	two	B	5	w
5	two	C	6	t



```
df.pivot(index='foo',  
          columns='bar',  
          values='baz')
```

bar	A	B	C
foo			
one	1	2	3
two	4	5	6



- Note!
  - It does not make sense to pivot every DataFrame, all of the datasets used in this course will have no need for a pivot table operation to use with machine learning.



# Pandas

- You should first go through this checklist **before** running a pivot():
  - What question are you trying to answer?
  - What would a dataframe that answers the question look like? Does it need a pivot()
  - What do you want the resulting pivot to look like?



# Pandas

- Pandas also comes with a `pivot_table` method that allows for an additional aggregation function to be called.
- This could alternatively be done with a `groupby()` method call as well.
- Let's explore both `.pivot()` and `pivot_table()` methods in pandas!



# Pivot Tables



# **Pandas Section Exercise - Overview**



# Pandas

- Let's test all your new pandas skills!
- Keep in mind:
  - Most questions can be solved in one or two lines of pandas code.
  - There could be multiple correct solutions.
  - Be careful not to run the cell above the expected output.



# **Pandas Section**

## **Exercise - Solutions**