


t-SNE

t-distributed Stochastic Neighbor Embedding

A visual, beginner-friendly handout for general audiences



Goal: keep nearby cases near each other, so patterns are easier to see.


One-sentence idea: t-SNE turns data with many features into a 2D or 3D map where similar cases tend to appear near one another.

No advanced mathematics required.

Beginner handout | t-SNE for visual exploration

1. Why do we need t-SNE?

Real datasets often describe each item using many measurements at once. Humans, however, are much better at seeing patterns on a flat page.



Goal: keep nearby cases near each other, so patterns are easier to see.

A person or object can have dozens, hundreds, or thousands of features. t-SNE compresses that information into a visual map.

What does “dimension” mean?

A dimension is simply one feature or variable. If a dataset records age, height, weight, and daily steps, it has four dimensions.

Think of a spreadsheet: Each column can be a dimension. A row is one person, sample, document, image, or other item.

Dimensionality reduction


Dimensionality reduction creates a simpler representation of complex data. t-SNE is mainly used to make a 2D or 3D visualization that helps us explore similarity and possible group structure.

- Input: many features per item.
- Output: usually two coordinates per item, so each item becomes a dot on a map.

t-SNE = t-distributed Stochastic Neighbor Embedding

2. The central idea: preserve neighborhoods

t-SNE asks a local question: “Which points are neighbors in the original data?” Then it tries to keep those neighbors close on the map.



t-SNE pays special attention to local neighbors.

The method gives high importance to nearby points. Distant points matter less than local neighbors.

A simple analogy

Imagine a large conference. You do not need an exact map of the whole city; you mainly want people with similar interests to stand near one another. t-SNE rearranges the room so that similar attendees are more likely to become neighbors.

Important: t-SNE is better at preserving local neighborhoods than global geometry. Two groups that look far apart on the page are not automatically “very different” in a precise numerical sense.

What the dots mean

- Each dot represents one observation.
- Nearby dots often represent observations with similar feature patterns.

t-SNE = t-distributed Stochastic Neighbor Embedding

3. How t-SNE works - without heavy math

The algorithm can be understood as four repeated ideas rather than a long formula.



The optimization tries to preserve local similarity, not every global distance.

1. Measure which items are similar in the original high-dimensional data.
2. Convert those similarities into probabilities: close neighbors receive larger probabilities.
3. Create a 2D map and measure which points are neighbors there.
4. Move the points so the 2D neighborhood probabilities resemble the original neighborhood probabilities as closely as possible.

Optimization in plain language: The map begins imperfectly. t-SNE repeatedly nudges points until local relationships are represented better.

Why “stochastic”?

The procedure can depend on random initialization or other random choices. Different runs can therefore produce maps that look rotated, reflected, or sometimes meaningfully different.

t-SNE = t-distributed Stochastic Neighbor Embedding

4. Why the “t” matters

The “t” in t-SNE refers to the Student t distribution used when comparing distances on the low-dimensional map.

The crowding problem

When many high-dimensional relationships are squeezed into only two dimensions, points can become crowded together. A simple 2D map does not have enough room to preserve every distance.

What the t distribution does: Its heavier tails allow moderately distant points to stay farther apart on the map. This gives local groups more room to form instead of collapsing into one crowded center.


Without enough room	With heavier tails
Different neighborhoods can be forced into the same small region.	Nearby points can stay close while unrelated points have more space between them.

You do not need to calculate the Student t distribution by hand to use t-SNE. The practical message is that it helps the 2D layout avoid excessive crowding.

t-SNE = t-distributed Stochastic Neighbor Embedding

5. A small example

Suppose each observation has 12 measured features. We cannot draw a 12-dimensional scatterplot, but t-SNE can create a 2D view.



Illustrative synthetic data. The colors are shown only to explain the example; t-SNE itself does not require class labels.

What can we reasonably say?

- Points inside the same compact region have similar high-dimensional neighborhoods.
- Separated regions may suggest different patterns worth investigating with the original variables or domain knowledge.

What we should not say: “The four groups are proven to be four real natural classes.” A t-SNE picture is exploratory evidence, not proof of true categories.

t-SNE is unsupervised

The algorithm usually receives only the feature values, not the group labels. Labels can be added afterward to help interpret the plot.

t-SNE = t-distributed Stochastic Neighbor Embedding

6. How to read a t-SNE plot


A good reading strategy is to start local, then verify any apparent pattern using the original data.

Usually meaningful

- Very close points often have similar feature patterns.
- Local neighborhoods or repeated small structures may be useful clues.

Often misleading

- The exact distance between far-away clusters.
- The absolute size, orientation, or position of a cluster on the page.



Rotation, reflection, and absolute axis values are not meaningful by themselves.

Axes are not ordinary variables: A t-SNE x-axis of 20 is not “twice” an x-axis value of 10 in any real-world sense. The coordinates mainly define the visualization layout.


Best habit

Treat the plot as a map for finding questions: Which points group together? Which observations are isolated? Then return to the original features to understand why.

t-SNE = t-distributed Stochastic Neighbor Embedding

7. Perplexity: the famous t-SNE setting

Perplexity controls the scale of neighborhood that t-SNE pays attention to. It is often described as a rough idea of the effective number of nearby neighbors.



Key idea: Different perplexity values can change the visual arrangement. Do not choose a value simply because it produces the “prettiest” clusters.

Practical guidance

- Try several reasonable values rather than trusting one run.
- Look for patterns that remain fairly stable across settings and random seeds.

Note: Perplexity must be smaller than the number of observations, and useful ranges depend on dataset size and structure.

t-SNE = t-distributed Stochastic Neighbor Embedding

8. Other settings and reproducibility

t-SNE is an optimization procedure, so settings and preprocessing can affect the final picture.

Random seed

Running t-SNE twice with different random seeds can produce different layouts. Similar local neighborhoods across runs are more reassuring than a pattern that appears only once.

Learning rate and iterations

The learning rate controls the size of optimization updates, while the number of iterations controls how long the map is refined. Modern software often provides reasonable automatic defaults, but convergence should still be checked.

Scaling the features

If one feature ranges from 0 to 1 and another from 0 to 1,000, the larger-scale feature may dominate distance calculations. Standardization or another suitable scaling method is often important.

For very high-dimensional data: A common workflow is to remove obvious noise and optionally use PCA first, then apply t-SNE to the reduced representation. This can improve speed and reduce noise.

Reproducible practice

- Record the software, random seed, perplexity, preprocessing, and other important settings.
- Save the code used to create the final visualization.

t-SNE = t-distributed Stochastic Neighbor Embedding

9. Strengths and limitations

What t-SNE is good at

- Making complex data visually explorable
- Showing local similarity and possible neighborhood structure
- Finding unusual points or subgroups to investigate
- Creating intuitive maps for images, text embeddings, biology, and other feature-rich data

What t-SNE is not good at

- Preserving exact long-range distances
- Proving that visible clusters are true categories
- Giving meaningful axis labels or directions
- Serving as a simple transformation for unseen new cases in classic t-SNE

The safest mental model: t-SNE is a visualization and exploration tool. It is excellent for generating hypotheses about local structure, but the original data should be used to test those hypotheses.

When another method may help

Use PCA when linear directions and explained variance matter. Consider UMAP when you want another nonlinear visualization method that often scales well and may preserve more broad structure. Method choice should follow the analysis goal.

t-SNE = t-distributed Stochastic Neighbor Embedding

10. A simple t-SNE workflow

Use this checklist when creating or reviewing a t-SNE visualization.

1. Prepare the data

Choose meaningful numeric features. Handle missing values and obvious data-quality problems.

2. Scale when appropriate

Standardize features when their units or ranges are not directly comparable.

3. Reduce noise if needed

For very large feature spaces, consider a preliminary PCA or another sensible filtering step.

4. Run t-SNE

Choose a random seed and try more than one perplexity value.

5. Repeat

Run several seeds/settings and look for stable local relationships.

6. Interpret carefully

Focus on nearby points and neighborhoods, not precise long-range distances.

7. Verify in the original data

Use feature summaries, domain knowledge, or formal analyses to explain apparent groups.

8. Report the settings

Document preprocessing, parameters, software, and seed.

Rule of thumb: If a scientific conclusion depends entirely on one t-SNE picture, the conclusion needs more evidence.

t-SNE = t-distributed Stochastic Neighbor Embedding

11. Quick recap, glossary, and references

Five things to remember

1. t-SNE turns many-dimensional data into a 2D or 3D visualization.
2. Its main goal is to preserve local neighborhoods.
3. Perplexity changes the neighborhood scale the method emphasizes.
4. Axes, rotations, cluster size, and far-away distances should not be overinterpreted.
5. Stable patterns should be checked against the original data and domain knowledge.

Mini glossary

Term	Plain-language meaning
Dimension	One measured feature or variable.
Embedding	A new representation of data, often in fewer dimensions.
Neighbor	An observation that is similar or close according to the feature space.
Perplexity	A t-SNE setting related to the effective neighborhood scale.
Unsupervised learning	Analysis performed without using target labels to guide the representation.
Visualization	A graphical representation used to explore and communicate patterns.

References and further reading

van der Maaten, L., & Hinton, G. (2008). Visualizing data using t-SNE. *Journal of Machine Learning Research*, 9, 2579-2605.

Wattenberg, M., Viégas, F., & Johnson, I. (2016). How to use t-SNE effectively. *Distill*.

scikit-learn documentation. t-SNE: T-distributed Stochastic Neighbor Embedding.

Final thought: Use t-SNE to discover questions, not to replace careful statistical or domain-based reasoning.